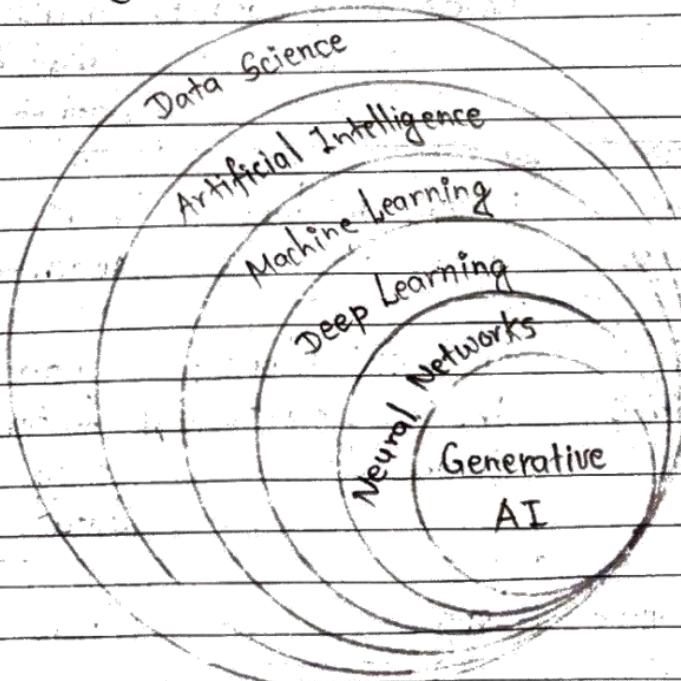


# g. Deep Learning and AI

## 1. Deep Learning

- Deep learning is a collection of statistical techniques of machine learning feature hierarchies that are actually based on artificial neural networks.
- While the term deep learning is vague , it bases on the idea of mimicking the brain by building algorithms that resemble biological neurons' functionality in the brain.



## Applications

- Computer Vision
- Natural Language Processing (NLP)

- Speech Recognition
- Health care
- Language translation
- Image detection
- Auto-pilot cars

## Deep Learning Vs Machine Learning

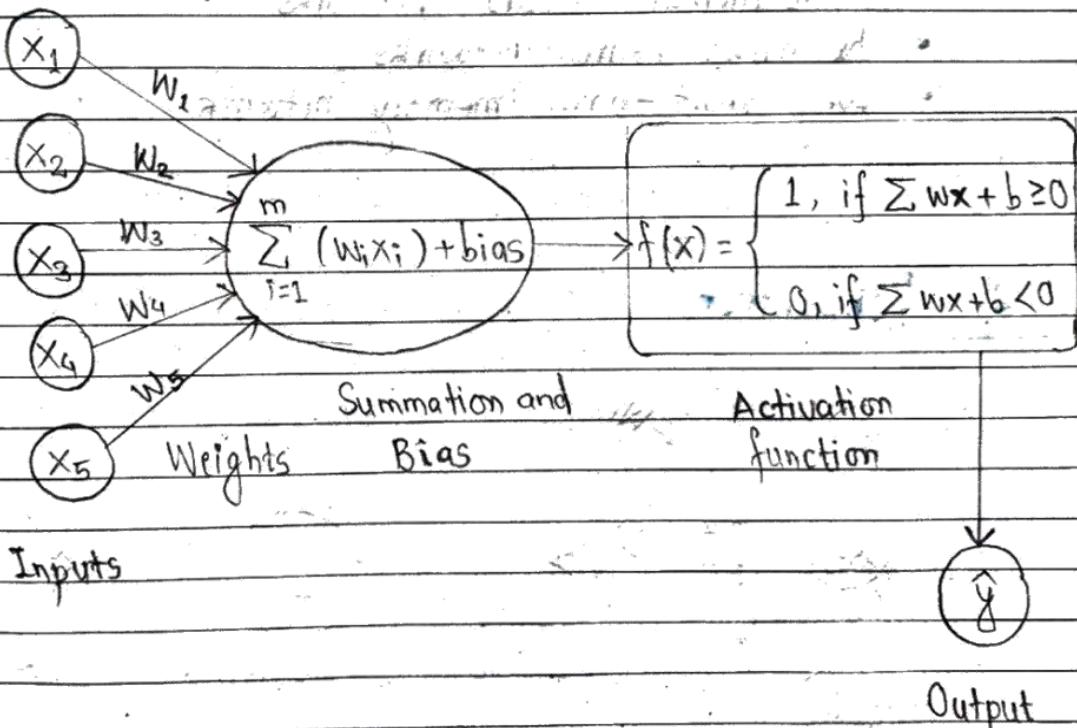
	Deep Learning	Machine Learning
Data	Needs a big dataset.	Performs well with a small to a medium dataset.
Hardware requirements	Requires machines with GPU.	Works with low-end machines.
Engineering peculiarities	Needs to understand the basic functionality of the data.	Understands the features and how they represent the data.
Training time	Long	Short
Processing time	A few hours or weeks.	A few seconds or hours.
Number of algorithms	Few	Many.

Data interpretation

Difficult

Some ML algorithms are easy and some are hard.

## Neural Network

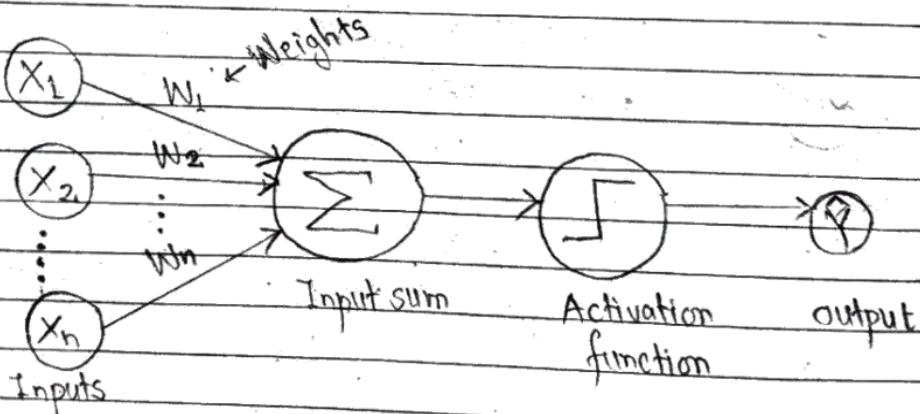


- 6 A neural network is a type of machine learning model by the structure and function of human brain.
- 6 It consists of interconnected units called neurons, which are organized into layers.
- 6 Structure : Input layer  $\rightarrow$  Hidden layer  $\rightarrow$  Output layer
- 6 Working : forward propagation and back propagation

## Types of Deep Learning Networks

- Perceptron (single layer)
- Feed forward networks
- Multi-layer perceptron (ANN)
- Radial based networks
- Convolutional neural networks
- Recurrent neural networks
- Long short-term memory network

### 2. Perceptron



6 A perceptron is one of the simplest type of artificial neural networks and serves as a fundamental building block for more complex neural network architectures.

### Working

1. Input features: The perceptron takes multiple input

features, each representing characteristics of the input data.

2. **Weights:** Each input feature is associated with a weight, which determines its influence on the output.
3. **Summation function:** The perceptron calculates the weighted sum of its input.
4. **Activation function:** The weighted sum is passed through an activation function (typically a step function) to produce the output. The output is usually binary (0 or 1).

### Working Practically

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from mlxtend.plotting import plot_decision_regions
```

```
dataset = pd.read_csv("Subscription.csv")
```

```
dataset.head(3)
```

```
# Check for null values
```

```
dataset.isnull().sum() # No null values.
```

```
# Visualize data
```

```
plt.figure(figsize=(9,3))
```

```
Sns.Scatterplot(x = "Age", y = "Salary", data = dataset,  
hue = "Purchase")  
plt.show()
```

```
# Separate input and output
```

```
x = dataset.iloc[:, :-1]
```

```
y = dataset["Purchase"]
```

```
# Train test split
```

```
from sklearn.model_selection import train-test-split
```

```
x-train, x-test, y-train, y-test = train-test-split(x, y,  
test-size = 0.2, random-state = 42)
```

```
# Perceptron model
```

```
from sklearn.linear_model import Perceptron
```

```
pr = Perceptron()
```

```
pr.fit(x-train, y-train)
```

```
pr.score(x-train, y-train), pr.score(x-test, y-test)
```

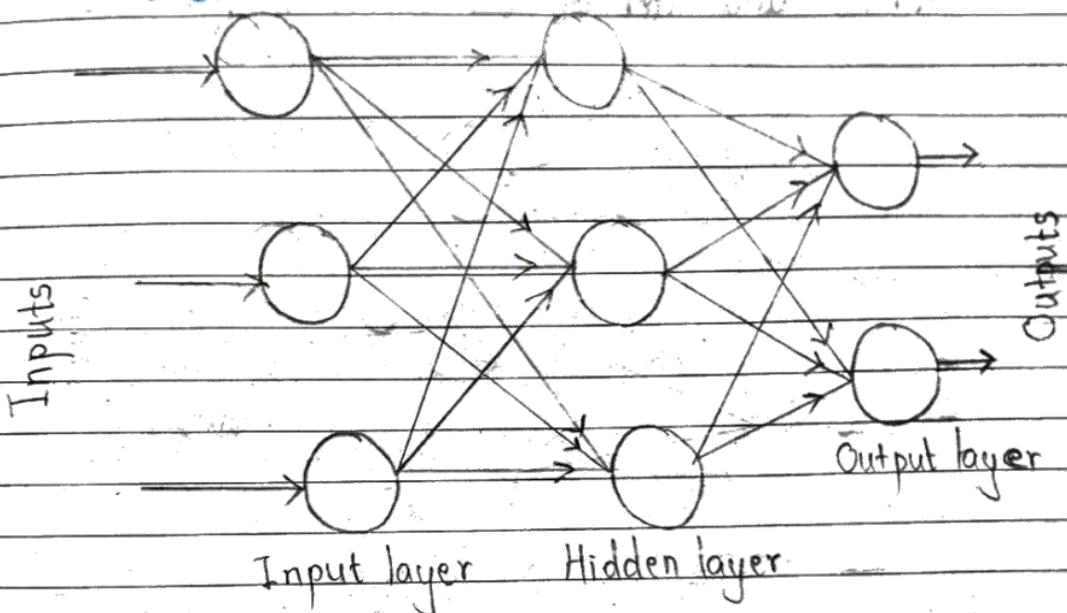
```
# Visualize decision region.
```

```
plot_decision-regions(x.to-numpy(), y.to-numpy(),  
clf = pr)
```

```
plt.show()
```

""This cannot provide high accuracy. So, instead of this multilayer perceptron is used""

### 3. Multilayer Perception.

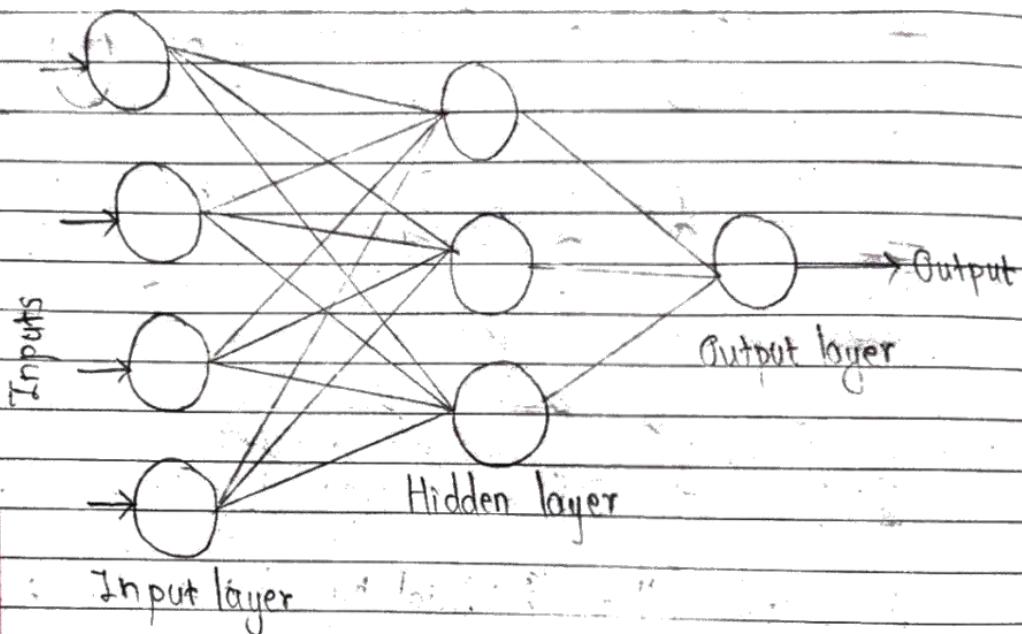


- It is also called Artificial Neural Networks (ANN)
- A multilayer perception (MLP) is a type of artificial neural network that consists of multiple layers of neurons, typically including an input layer, one or more hidden layers and one output layer.
- Working is based on forward propagation and back-propagation.

#### Forward Propagation

- Data is passed through the network from the input layer to the output layer.

- Each neuron computes a weighted sum of its inputs and applies an activation function.



### Back - Propagation

- Backpropagation is a technique of adjustment of weights.
- The error between the predicted output and actual output is calculated. The error is propagated back through the network and weights are adjusted to minimize the error.
- Uses gradient descent technique (formula).  
i.e.  $W_{new} = W_{old} - \lambda \left( \frac{dE}{dW} \right)$   
where  $\lambda$  is learning rate

## Working Practically

```
import pandas as pd
```

```
dataset = pd.read_csv("ChurnModelling.csv")  
dataset.head(3)
```

```
# Drop columns with object (string) type data  
dataset = dataset.select_dtypes(exclude=["object"])  
dataset.head(3)
```

```
# Check for null values
```

```
dataset.isnull().sum() # No null values
```

```
# Separate input and output data
```

```
input_data = dataset.iloc[:, :-1]  
output_data = dataset.iloc[:, -1]
```

```
# Scaling the input data
```

```
from sklearn.preprocessing import StandardScaler
```

```
ss = StandardScaler()
```

```
input_data = pd.DataFrame(ss.fit_transform(input_data),  
columns=input_data.columns)
```

```
# Train test split
```

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(input_data,  
output_data, test_size=0.2, random_state=42)
```

```
# Create an artificial neural networks  
import tensorflow  
import keras.layers import Dense  
import keras.models import Sequential  
  
ann = Sequential() # Object initialization  
  
# Check the shape of training data  
print(x-train.shape, y-train.shape) # Ensure shape are  
# correct  
  
# Build different layers.  
ann.add(Dense(8, input_dim=x-train.shape[1],  
activation="relu")) # first hidden layer  
# first parameter 8 is the number of nodes in the  
# layer  
ann.add(Dense(6, activation="relu")) # Second hidden  
# layer  
# Input dimension is needed only for first hidden layer  
ann.add(Dense(4, activation="relu")) # Third hidden  
# layer  
ann.add(Dense(2, activation="relu")) # fourth hidden  
# layer  
ann.add(Dense(1, activation = "sigmoid")) # Output  
# layer  
# Sigmoid is used because our output is binary  
  
# Compile the model  
ann.compile(optimizer="adam", loss="binary_crossentropy",  
metrics=["accuracy"])
```

# Train the model.

ann.fit(x\_train, y\_train, batch\_size=100, epochs=50)

# Using batches helps in efficient computation and  
# faster convergence.

# Batches allows the model to update its weights

# more frequently than if it were using the entire  
# dataset

# Multiple epochs are used to train the model,  
# allowing it to learn and adjust its weights  
# iteratively.

# Check the accuracy of our model.

from sklearn.metrics import accuracy\_score

# We need prediction value of testing data.

ann.predict(x\_test) # The data is not in binary form.

# Convert the data in binary form

prediction = ann.predict(x\_test)

binary\_prediction = []

for i in prediction:

if i[0] > 0.5:

binary\_prediction.append(1)

else:

binary\_prediction.append(0)

# View prediction

binary\_prediction

accuracy\_score(y\_test, binary\_prediction) \* 100

# Accuracy of testing data

# For accuracy of training data

# We need prediction value of training data in  
# binary format

prediction = ann.predict(x-train)

binary-prediction = [ ]

for i in prediction :

if i[0] > 0.5 :

binary-prediction.append(1)

else:

binary-prediction.append(0)

# View prediction

binary-prediction

accuracy-score(y-train, binary-prediction) \* 100

# Accuracy of training data.

# To make the new predictions

Import numpy as np

new-data = ss.fit\_transform([[1, 15634602, 619, 42,  
2, 0.08, 1, 1, 1, 101848.88]])

new-prediction = ann.predict(new-data)

if new-prediction > 0.5 :

binary-form = 1

else:

binary-form = 0

# View output

binary-form

## 4. Activation Functions

- ↳ An activation function decides whether a neuron should be activated or not. This means that it will decide whether the neuron's input to the network is important or not in the process of prediction using simpler mathematical operations.
- ↳ The activation function is categorized into three main parts:
  1. Binary step function
  2. Linear activation function
  3. Non-linear activation function.

### 1. Binary Step Function

- ↳ Binary step function depends on a threshold value that decides whether a neuron should be activated or not.

- ↳ Mathematically, binary step function is:

$$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$$

where,

$$f(x) = \hat{y} = x_1 w_1 + x_2 w_2 + \dots + x_n w_n + b$$

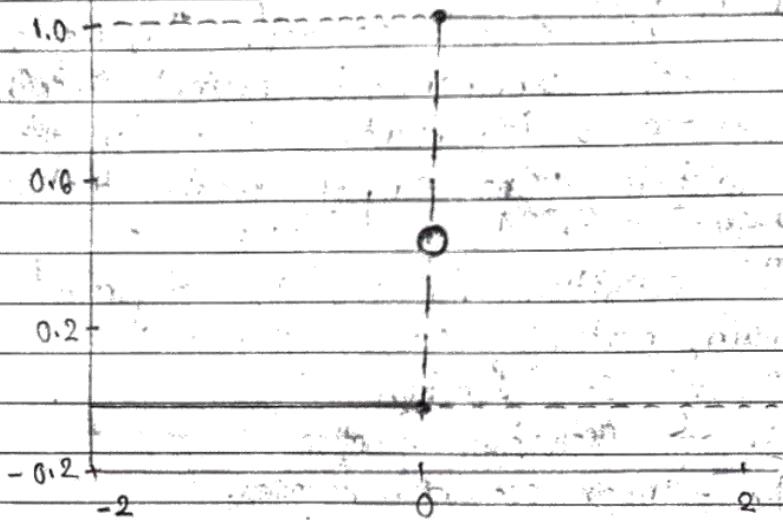
where,

$\hat{y}$  is predicted output

$x_1, x_2, \dots, x_n$  are input values

$w_1, w_2, \dots, w_n$  are weightages.

## Binary Step Activation Function



## 2. Linear Activation Function

- 6. The output of the functions is not restricted in between any range.
- 6. Its range is specified from  $-\infty$  to  $\infty$ .

6. Mathematically, linear activation function is:

$$f(x) = x$$

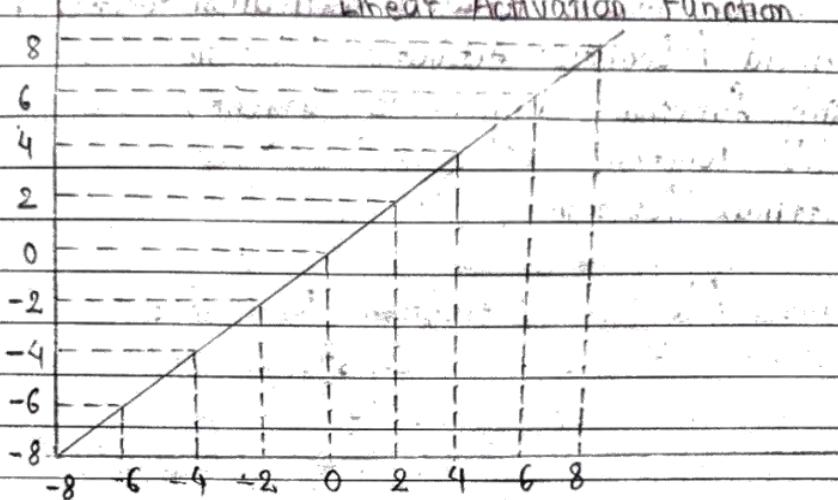
i.e.

$$\hat{y} = x$$

where,

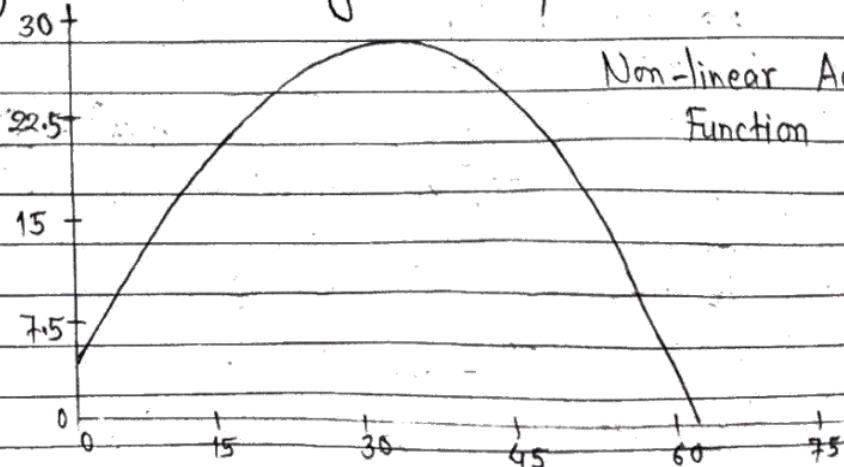
$\hat{y}$  is predicted output and  
 $x$  is input.

### 2. Linear Activation Function



### 3. Non-Linear Activation Functions

- ↳ These are one of the mostly widely used activation functions.
- ↳ It helps the model in generalizing and adapting any sort of data in order to perform correct differentiation among the output.

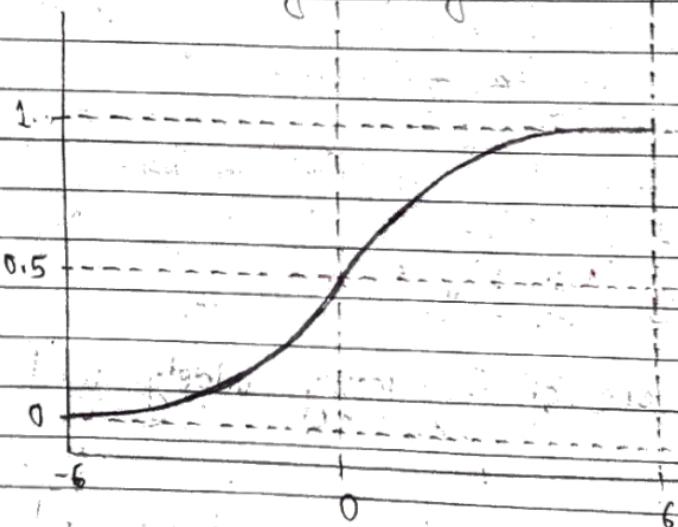


## 6 Non-linear Neural Networks Activation functions:

- i. Sigmoid / Logistic Activation Function
- ii. Tanh function (Hyperbolic Tangent)
- iii. ReLU Function
- iv. Softmax function

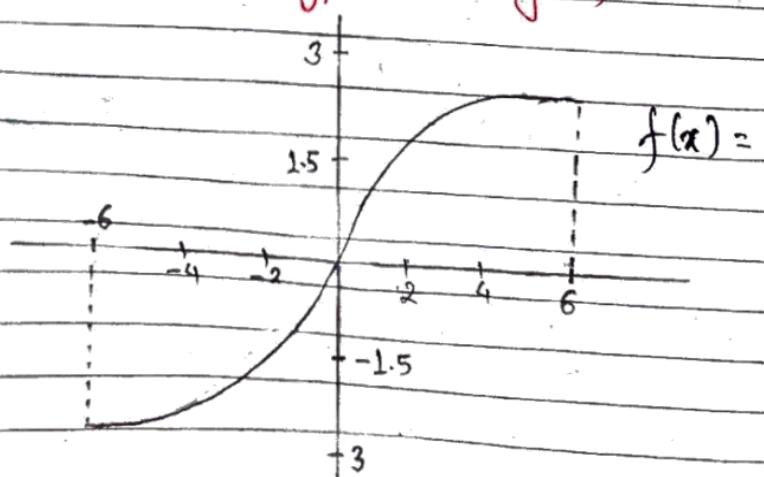
### i. Sigmoid / Logistic Activation Function :

Sigmoid / Logistic Activation function



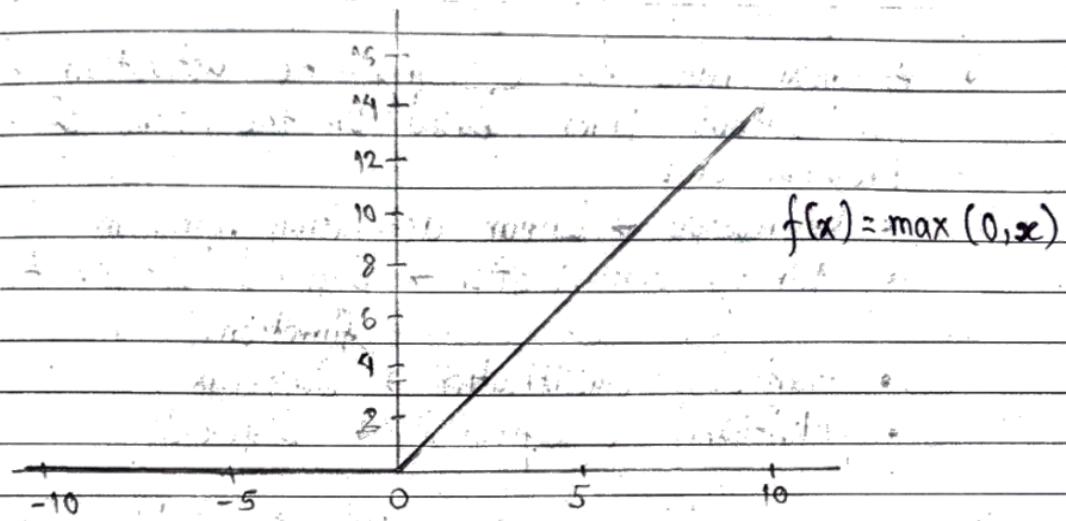
$$f(x) = \frac{1}{1+e^{-x}}$$

### ii. Tanh Function (Hyperbolic Tangent) :

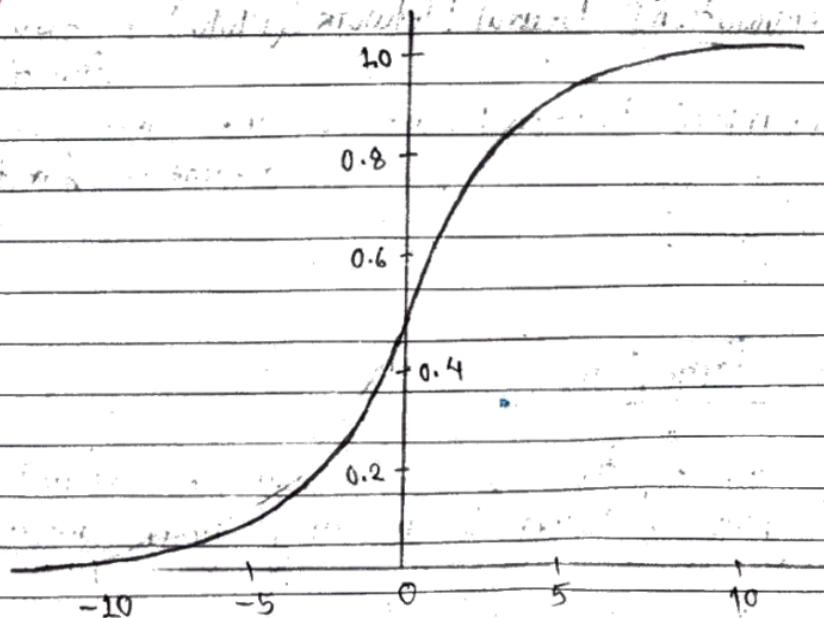


$$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

### III. ReLU Function:



### IV. Softmax Function:



$$\text{softmax}(z_i) = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$$

## Rules for Choosing Right Activation Function

- ↳ A few rules for choosing the activation function for output layer based on the type of prediction problem are:
  - Regression → Linear activation function
  - Binary classification → Sigmoid / Logistic Activation function.
  - Multi-class classification → Softmax
  - Multilabel classification → Sigmoid
- ↳ The activation function used in hidden layers is typically chosen based on the type of neural network architecture:
  - Convolutional Neural Network (CNN) : ReLU activation function
  - Recurrent Neural Network : Tanh and/or Sigmoid activation function

## 5. Loss Functions

- ↳ The loss function is a method of evaluating how well your algorithm is modeling your dataset.
- ↳ It is the mathematical function of the parameters of the machine learning algorithm.
- ↳ Remember that cost function is average of loss function of all data in dataset.

## Types of Loss Functions

### ■ Regression :

- MSE (Mean Squared Error)
- MAE (Mean Absolute Error)
- Huber loss

### ■ Classification :

- Binary cross-entropy / log loss
- Categorical cross-entropy

### ■ Auto Encoder:

- KL Divergence.

### ■ GAN :

- Discriminator loss
- Minmax GAN loss

### ■ Object Detection :

- Focal loss

### ■ Word Embeddings :

- Triplet loss

### ■ Regression Loss

#### 1. Mean Squared Error / Squared loss / L2 loss:

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (Y_i - \hat{Y}_i)^2$$

## 2. Mean Absolute Error / L1 Loss:

$$\text{MAE} = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$

## 3. Huber Loss:

$$\text{Huber} = \frac{1}{n} \sum_{i=1}^n \frac{1}{2} (y_i - \hat{y}_i)^2, |y_i - \hat{y}_i| \leq \delta$$

$$\text{Huber} = \frac{1}{n} \sum_{i=1}^n \delta (|y_i - \hat{y}_i| - \frac{1}{2}\delta), |y_i - \hat{y}_i| > \delta$$

where,

$n$  = total number of data points

$y$  = the actual value of the data point. Also known as true value.

$\hat{y}$  = the predicted value of the data point. This value is returned by the model.

$\delta$  = defines the point where the Huber loss function transitions from a quadratic to linear.

### Remember:

- MSE is used when there are no outliers.
- MAE is used when there are outliers but the function is not differentiable
- Huber loss is used when you have around 30% outliers in dataset.

## Classification Loss

### 1. Binary Cross Entropy / Log Loss:

$$\text{Log loss} = -\frac{1}{N} \sum_{i=1}^N y_i \log \hat{y}_i + (1-y_i) \log (1-\hat{y}_i)$$

Log loss works when we have binary output.

So,

When output = 0 ;

$$\text{log loss} = -\frac{1}{N} \sum_{i=1}^N \log (1-\hat{y}_i)$$

When output = 1 ;

$$\text{log loss} = -\frac{1}{N} \sum_{i=1}^N \log \hat{y}_i$$

### 2. Categorical Cross Entropy:

$$\text{loss} = - \sum_{j=1}^k y_j \log (\hat{y}_j)$$

Where,

$k$  is number of classes in the data.

- If target column has one hot encode to classes like 001, 010, 100 then use categorical cross entropy.
- If the target column has numerical encoding to classes like 1, 2, 3, 4, ..., N then use sparse categorical cross-entropy.

## 6. Optimizer in Neural Network

- 6 Optimizers are algorithms or methods used to change the attributes of your neural network such as weights and learning rate in order to reduce the losses.
- 6 Used when we have lots of local minimums and a global minimum where algo can stick on local minimum.
- 6 Some Optimizers are:
  - Gradient descent
  - Stochastic gradient descent
  - Stochastic gradient descent with momentum
  - Mini-batch gradient descent
  - Adagrad
  - RMSProp
  - AdaDelta
  - Adam

## 7. Improve the Performance of Neural Network

### 6 Hyperparameters Tuning:

- No. of hidden layers
- No. of neurons per layer
- Learning rate
- Optimizer
- Batch size

- Activation function
- Epochs

#### ↳ Solving Vanishing / Exploding Gradient :

- Optimizer manipulation
- Changing initialization weight
- Batch normalization.

#### ↳ Use Big Data :

- Use transform learning to increase size of data.

#### ↳ Overcome Slow Training :

- Can be achieved by hyperparameter tuning.
- Solve vanishing/exploding gradient.

#### ↳ Control Overfitting :

- L1 and L2 regularization
- Early stopping
- Batch normalization.

### 8. Overfitting in Deep Learning

↳ Overfitting is a common explanation for the poor performance of a predictive model.

↳ Overfitting refers to an unwanted behaviour of a machine learning algorithm used for predictive modeling.

## Minimization of Overfitting

- Cross validation
- Train with more data
- Remove features
- Early stopping
- Regularization (L1 and L2)
- Ensembling
- Hyperparameter tuning

### Early Stopping

- ↳ Early stopping is a regularization technique in machine learning to prevent overfitting.
- ↳ It involves monitoring the model's performance on a validation set during training and stopping the training process when the performance on the validation set starts to degrade.

### Regularization

- ↳ Regularization is a technique used in machine learning to prevent overfitting by adding a penalty to the loss function.
- ↳ This penalty discourages the model from becoming too complex and helps it generalize better to new, unseen dataset.

## Working Practically

```
import pandas as pd  
import matplotlib.pyplot as plt
```

```
dataset = pd.read_csv("ChurnModeling.csv")  
dataset.head(3)
```

```
# Drop columns with object (string) type data  
dataset = dataset.select_dtypes(exclude=["object"])  
dataset.head(3)
```

```
# Check for null values  
dataset.isnull().sum() # No null values
```

```
# Separate input and output data  
input_data = dataset.iloc[:, :-1]  
output_data = dataset.iloc[:, -1]
```

```
# Scaling the input data  
from sklearn.preprocessing import StandardScaler
```

```
ss = StandardScaler()  
input_data = pd.DataFrame(ss.fit_transform(input_data),  
                          columns = input_data.columns)
```

```
# Train test split  
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(input_data,  
                                                    output_data, test_size = 0.2, random_state = 12)
```

# Create an artificial neural network

```
import tensorflow
```

```
from keras.layers import Dense
```

```
from keras.models import Sequential
```

```
ann = Sequential() # Object initialization
```

# Check the shape of training data

```
print(x_train.shape, y_train.shape) # Ensure shapes are  
# correct
```

# Build different layers

```
ann.add(Dense(8, input_dim=x_train.shape[1],  
activation="relu")) # first hidden layer
```

# first parameter 8 is the number of nodes in the  
# layer

```
ann.add(Dense(6, activation="relu")) # second hidden  
# layer
```

```
ann.add(Dense(4, activation="relu")) # third hidden  
# layer
```

```
ann.add(Dense(2, activation="relu")) # fourth hidden  
# layer
```

```
ann.add(Dense(1, activation="sigmoid")) # Output  
# layer
```

# Compile the model

```
ann.compile(optimizer="adam", loss="binary_crossentropy",  
metrics=["accuracy"])
```

# Train the model

```
ann.fit(x_train, y_train, batch_size=100, epochs=50,
```

validation\_data = (x-test, y-test)

# Validation data is for testing accuracy.

# View training history  
ann.history.history

# The history attribute of a keras model contain  
# the training history, including metrics like loss and  
# accuracy for each epoch.

# To get all the keys of history  
ann.history.history.keys()

# To see specific values of keys

ann.history.history["accuracy"] # for accuracy

# To get training and testing accuracy of each epoch

train-accuracy = ann.history.history["accuracy"]

test-accuracy = ann.history.history["val-accuracy"]

# Visualize the accuracy

# Get the lengths of accuracies for plotting  
len(test-accuracy), len(train-accuracy)

# Draw graph

plt.plot([i+1 for i in range(len(test-accuracy))],  
test-accuracy, color="blue")

plt.plot([i+1 for i in range(len(train-accuracy))],  
train-accuracy, color="red")

plt.show()

# Check the accuracy of our model

from sklearn.metrics import accuracy\_score

# Predictions of x-train

predict\_x\_train = ann.predict(x\_train)

x\_train\_predictions = []

for i in predict\_x\_train:

if i[0] > 0.5 :

x\_train\_predictions.append(1)

else:

x\_train\_predictions.append(0)

# View predictions

x\_train\_predictions

# Predictions of x-test

predict\_x\_test = ann.predict(x\_test)

x\_test\_predictions = []

for i in predict\_x\_test:

if i[0] > 0.5 :

x\_test\_predictions.append(1)

else:

x\_test\_predictions.append(0)

# View predictions

x\_test\_predictions

# Get accuracy

train\_data\_accuracy = accuracy\_score(y\_train, x\_train\_predictions) \* 100

test\_data\_accuracy = accuracy\_score(y\_test, x\_test\_predictions) \* 100

train\_data\_accuracy, test\_data\_accuracy.

"" If the training accuracy is more than testing accuracy, the model is overfitted.

To avoid overfitting we use various techniques like Early Stopping and Regularization.

These methods help in minimizing the gaps between the accuracies. ""

## Early Stopping :

from keras.callbacks import EarlyStopping

# Train the model

ann.fit (x-train, y-train, batch-size = 100, epochs = 50,  
validation-data = (x-test, y-test), callbacks =  
EarlyStopping ())

# Early Stopping stops training if epoch is going towards  
# overfitting

# To get training and testing accuracy

train-accuracy = ann.history.history ["accuracy"]  
test-accuracy = ann.history.history ["val-accuracy"]

# Visualize the accuracy

plt.plot ([i+1 for i in range (len (test-accuracy))],  
test-accuracy, color = "blue")

plt.plot ([i+1 for i in range (len (train-accuracy))],  
train-accuracy, color = "red")

plt.show()

## Regularization :

```
from keras.regularizers import L2
```

```
# Build different layers
```

```
ann.add(Dense(8, input_dim=x_train.shape[1],  
activation="relu",  
kernel_regularizer=L2(l2=0.01)))
```

```
# kernel-regularizer stops the training if it is going  
# towards overfitting using given regularization  
# technique.
```

```
ann.add(Dense(6, activation="relu", kernel_regularizer=  
L2(l2=0.01)))
```

```
# Kernel-regularizer can be used in different layers  
# differently if you want or you can simply skip as  
# below
```

```
ann.add(Dense(4, activation="relu"))
```

```
ann.add(Dense(2, activation="relu"))
```

```
ann.add(Dense(1, activation="sigmoid"))
```

```
# Compile the model
```

```
ann.compile(optimizer="adam", loss="binary-  
crossentropy", metrics=["accuracy"])
```

```
# Train the model
```

```
ann.fit(x_train, y_train, batch_size=200, epochs=50,  
validation_data=(x_test, y_test))
```

```
# To get training and testing accuracy.
```

```
train_accuracy = ann.history.history["accuracy"].
```

```
test_accuracy = ann.history.history["val_accuracy"]
```

```
train_accuracy, test_accuracy
```

#Visualize the accuracy

```
plt.plot ([i+1 for i in range (len (test_accuracy))],  
         test_accuracy , color = "blue")  
plt.plot ([i+1 for i in range (len (train_accuracy))],  
         train_accuracy , color = "red")  
plt.show ()
```

## 9. Batch Normalization

- Batch normalization is a technique used to improve the training of deep neuron networks by normalizing the inputs of each layer
- This helps to stabilize and accelerate the training process.

### Working Practically

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
dataset = pd.read_csv ("ChurnModelling.csv")  
dataset.head (3)
```

#Drop columns with object (string) type data

```
dataset = dataset.select_dtypes (exclude = ["object"] )  
dataset.head (3)
```

```
# Check for null values
```

```
dataset.isnull().sum() # No null values
```

```
# Separate input and output data
```

```
input_data = dataset.iloc[:, :-1]
```

```
output_data = dataset.iloc[:, -1]
```

```
# Scaling the input data
```

```
from sklearn.preprocessing import StandardScaler
```

```
ss = StandardScaler()
```

```
input_data = pd.DataFrame(ss.fit_transform(input_data),  
columns = input_data.columns)
```

```
# Train test split
```

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(  
input_data, output_data, test_size=0.2,  
random_state=42)
```

```
# Create an artificial neural network
```

```
import tensorflow
```

```
from keras.layers import Dense
```

```
from keras.models import Sequential
```

```
from keras.callbacks import EarlyStopping
```

```
from keras.regularizers import L2
```

```
ann = Sequential() # Object initialization
```

```
# Build different layers.
```

```
ann.add(Dense(8, input_dim=x_train.shape[1],
activation="relu", kernel_regularizer=L2(l2=0.01))) # first hidden
```

# layer.

```
ann.add(Dense(6, activation="relu", kernel_regularizer=L2(l2=0.01)))
```

```
ann.add(Dense(4, activation="relu", kernel_regularizer=L2(l2=0.01)))
```

```
ann.add(Dense(2, activation="relu", kernel_regularizer=L2(l2=0.01)))
```

```
ann.add(Dense(1, activation="sigmoid"))
```

# Compile the model

```
ann.compile(optimizer="adam", loss="binary_crossentropy",
metrics=["accuracy"])
```

# Train the model

```
ann.fit(x_train, y_train, batch_size=100, epochs=50,
validation_data=(x_test, y_test), callbacks=[EarlyStopping()])
```

# Training and testing accuracy of each epoch

```
train_accuracy = ann.history.history["accuracy"]
```

```
test_accuracy = ann.history.history["val_accuracy"]
```

train\_accuracy, test\_accuracy

# Visualize the accuracy

```
plt.plot([i+1 for i in range(len(test_accuracy))], test_accuracy, color="blue")
```

```
plt.plot([i+1 for i in range(len(train_accuracy))], train_accuracy, color="red")
```

plt.show()

# Check the accuracy of our model  
from sklearn.metrics import accuracy\_score

# Predictions of x-train

predict\_x\_train = ann.predict(x\_train)

x\_train\_predictions = []

for i in predict\_x\_train:

if i[0] > 0.5:

x\_train\_predictions.append(1)

else:

x\_train\_predictions.append(0)

# View predictions

x\_train\_predictions

# Predictions of x-test

predict\_x\_test = ann.predict(x\_test)

x\_test\_predictions = []

for i in predict\_x\_test:

if i[0] > 0.5:

x\_test\_predictions.append(1)

else:

x\_test\_predictions.append(0)

# View predictions

x\_test\_predictions

# Get accuracy

train\_data\_accuracy = accuracy\_score(y\_train, x\_train\_predictions) \* 100

test\_data\_accuracy = accuracy\_score(y\_test, x\_test\_predictions)

predictions) \* 100  
train-data-accuracy, test-data-accuracy

"" If the training accuracy is more than testing accuracy, the model is overfitted.

To avoid overfitting we use Batch Normalization to minimize the gaps between the accuracies. """

### Batch Normalization:

```
from keras.layers import BatchNormalization
```

```
# Build different layers
```

```
ann.add(Dense(8, input_dim=x_train.shape[1],  
activation = "relu",  
kernel_regularization = L2(12=0.01)))
```

```
ann.add(BatchNormalization()) # Applying batch  
# normalization on first hidden layer output
```

```
ann.add(Dense(6, activation = "relu",  
kernel_regularization = L2(12=0.01)))
```

```
ann.add(BatchNormalization()) # Applying batch  
# normalization on second hidden layer output.
```

```
ann.add(Dense(4, activation = "relu",  
kernel_regularization = L2(12=0.01)))
```

```
ann.add(BatchNormalization()) # Applying batch  
# normalization on third hidden layer output.
```

```
ann.add(Dense(2, activation = "relu",  
kernel_regularization = L2(12=0.01)))
```

```
ann.add(BatchNormalization()) # Applying batch  
# normalization on fourth hidden layer.
```

```
ann.add(Dense(1, activation = "sigmoid"))
```

# Batch normalization is not required in output  
 # layer

# Compile the model

ann.compile(optimizer = "adam", loss = "binary-crossentropy", metrics = ["accuracy"])

# Train the model

ann.fit(x\_train, y\_train, batch\_size = 100, epochs = 50,  
 validation\_data = (x\_test, y\_test),  
 callbacks = EarlyStopping())

# Training and testing accuracy of each epoch.

train\_accuracy = ann.history.history["accuracy"]

test\_accuracy = ann.history.history["val\_accuracy"]

train\_accuracy, test\_accuracy

# Visualize the accuracy

plt.plot([i+1 for i in range(len(test\_accuracy))],  
 test\_accuracy, color = "blue")

plt.plot([i+1 for i in range(len(train\_accuracy))],  
 train\_accuracy, color = "red")

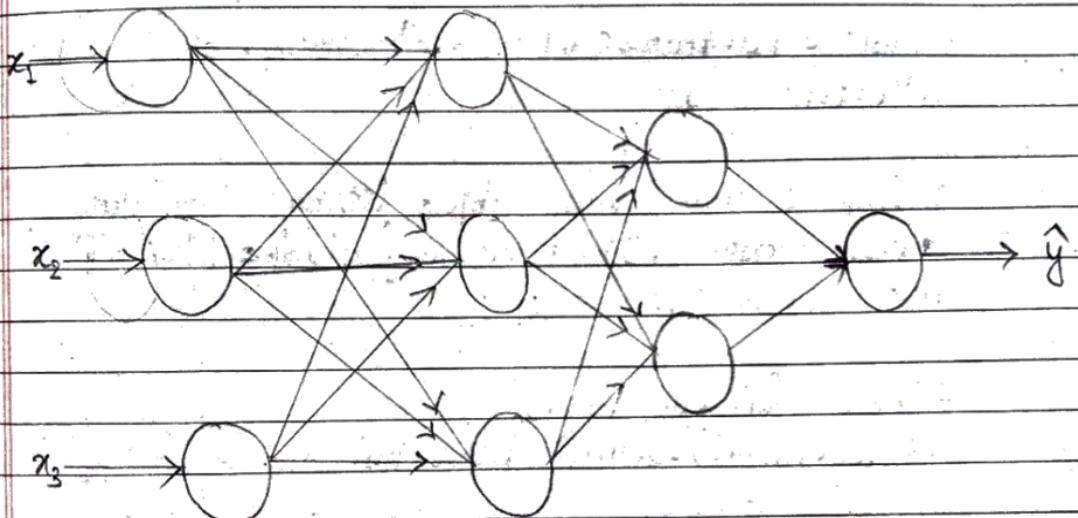
plt.show()

## 10. Dropout Layer

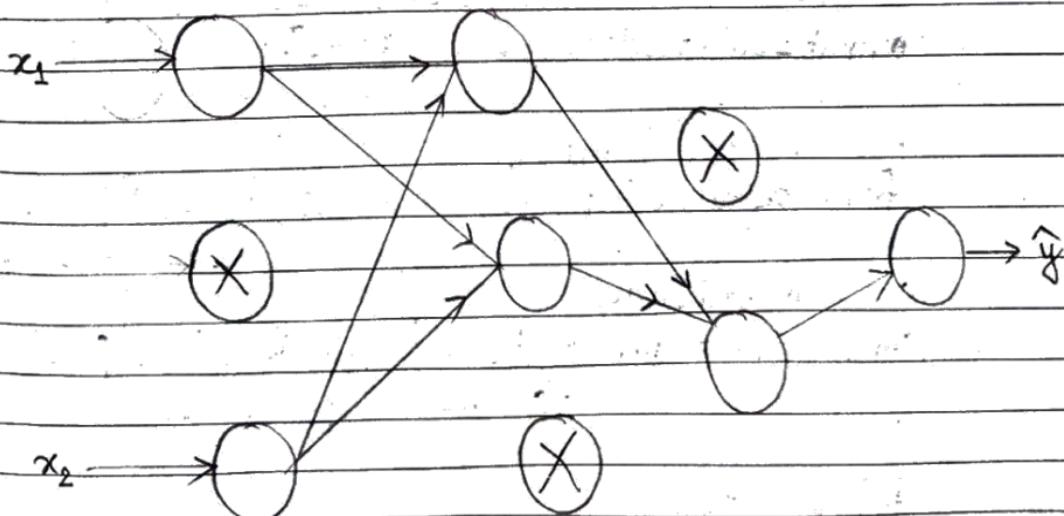
- ↳ All the forward and backward connections with a dropped node are temporarily removed, thus creating a new network architecture out of

the parent network.

6 The nodes are dropped by a dropout probability of  $p$ .



(a) Standard Neural Net



(b) After applying dropout layer

## Working Practically

```
import pandas as pd  
import matplotlib.pyplot as plt
```

```
dataset = pd.read_csv("Churn Modeling.csv")  
dataset.head(3)
```

```
# Drop columns with object (string) type data  
dataset = dataset.select_dtypes(exclude=["object"])  
dataset.head(3)
```

```
# Check for null values  
dataset.isnull().sum() # No null values
```

```
# Separate input and output data.  
input_data = dataset.iloc[:, :-1]  
output_data = dataset.iloc[:, -1]
```

```
# Scaling the input data  
from sklearn.preprocessing import StandardScaler
```

```
ss = StandardScaler()
```

```
input_data = pd.DataFrame(ss.fit_transform(input_data),  
                           columns=input_data.columns)
```

```
# Train test split
```

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(  
    input_data, output_data, test_size=0.2,
```

random\_state = 42.)

# Create an artificial neural network

import tensorflow

from keras.layers import Dense, BatchNormalization

from keras.models import Sequential

from keras.callbacks import EarlyStopping

from keras.regularizers import L2

ann = Sequential() # Object initialization

# Build different layers

ann.add(Dense(8, input\_dim=x\_train.shape[1], activation = "relu", kernel\_regularizer = L2(l2=0.01)))

ann.add(BatchNormalization())

ann.add(Dense(6, activation = "relu", kernel\_regularizer = L2(l2=0.01)))

ann.add(BatchNormalization())

ann.add(Dense(4, activation = "relu", kernel\_regularizer = L2(l2=0.01)))

ann.add(BatchNormalization())

ann.add(Dense(2, activation = "relu", kernel\_regularizer = L2(l2=0.01)))

ann.add(BatchNormalization())

ann.add(Dense(1, activation = "sigmoid"))

# Compile the model

ann.compile(optimizer = "adam", loss = "binary\_crossentropy", metrics = ["accuracy"])

# Train the model

```
ann.fit(x-train, y-train, batch-size=100, epochs=50, validation-data=(x-test, y-test),  
       callbacks=EarlyStopping())
```

# Training and testing accuracy of each epoch

```
train-accuracy = ann.history.history["accuracy"]  
test-accuracy = ann.history.history["val-accuracy"]  
train-accuracy, test-accuracy
```

# Visualize the accuracy

```
plt.plot([i+1 for i in range(len(test-accuracy))],  
         test-accuracy, color="blue")  
plt.plot([i+1 for i in range(len(train-accuracy))],  
         train-accuracy, color="red")  
plt.show()
```

# Check the accuracy of our model

```
from sklearn.metrics import accuracy-score
```

# Predictions of x-train

```
predict-x-train = ann.predict(x-train)
```

```
x-train-predictions = []
```

```
for i in predict-x-train:
```

```
    if i[0] > 0.5:
```

```
        x-train-predictions.append(1)
```

```
    else:
```

```
        x-train-predictions.append(0)
```

# View predictions

```
x-train-predictions
```

# Prediction of x-test

predict-x-test = ann.predict(x-test)

x-test-predictions = []

for i in predict-x-test:

if i[0] > 0.5:

x-test-predictions.append(1)

else:

x-test-predictions.append(0)

# View predictions

x-test-predictions

# Get accuracy

train-data-accuracy = accuracy-score(y-train,  
x-train-predictions) \* 100

test-data-accuracy = accuracy-score(y-test,  
x-test-predictions) \* 100

train-data-accuracy, test-data-accuracy

" If the training accuracy is more than testing accuracy then the model is overfitted.

To avoid overfitting, we can use Dropout layer to minimize the gap between the accuracies."

### Dropout Layer:

from keras.layers import Dropout

# Build different layers

ann.add(BatchNormalization())

ann.add(Dense(8, input\_dim=x-train.shape[1],  
activation = "relu"))

```
kernel_regularizer = L2 (l2=0.01)))  
ann.add (Dropout (0.3)) # Dropout layer with 30%  
# to first layer  
ann.add (Batch Normalization ())  
ann.add (Dense (6, activation = "relu",  
    kernel_regularizer = L2 (l2=0.01)))  
ann.add (Dropout (0.3)) # Dropout layer to second  
# layer  
ann.add (Batch Normalization ())  
ann.add (Dense (4, activation = "relu",  
    kernel_regularizer = L2 (l2=0.01)))  
ann.add (Dropout (0.3)) # Dropout layer to third  
# layer  
ann.add (Batch Normalization ())  
ann.add (Dense (2, activation = "relu",  
    kernel_regularizer = L2 (l2=0.01)))  
ann.add (Dropout (0.3)) # Dropout layer to fourth  
# layer  
ann.add (Batch Normalization ())  
ann.add (1, activation = "sigmoid"))  
  
# Compile the model  
ann.compile (optimizer = "adam", loss = "binary-  
crossentropy", metrics = ["accuracy"])  
  
# Train the model  
ann.fit (x_train, y_train, batch_size = 100, epochs = 50,  
validation_data = (x_test, y_test), callbacks =  
EarlyStopping ())  
  
# Training and testing accuracy of each epoch
```

train-accuracy = ann.history.history["accuracy"]

test-accuracy = ann.history.history["val-accuracy"]

train-accuracy, test-accuracy

#Visualize the accuracy

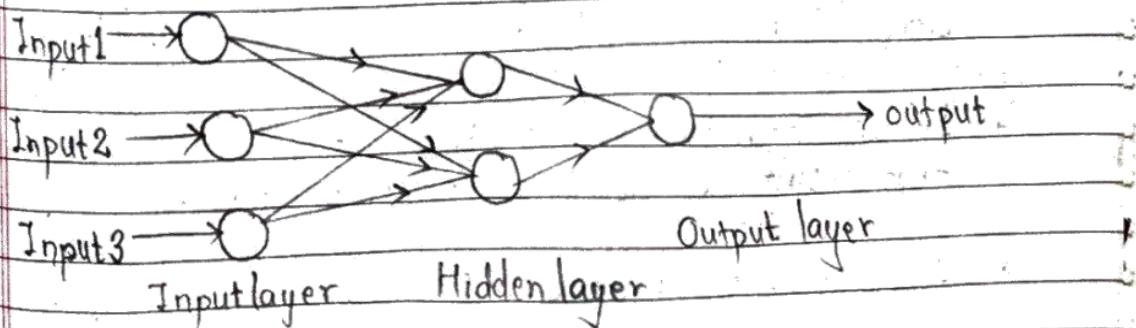
```
plt.plot([i+1 for i in range(len(test-accuracy))],  
         test-accuracy, color = "blue")
```

```
plt.plot([i+1 for i in range(len(train-accuracy))],  
         train-accuracy, color = "red")
```

```
plt.show()
```

## 11. Vanishing Gradient Problem

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods and backpropagation.
- In such methods, during each iteration of training each of the neural network's weights receives an update proportional to the partial derivative of the error function with respect to the current weight.



## Solutions for Vanishing Gradient Problem

1. Proper weight initialization
2. Using non-saturating activation functions (relu)
3. Batch normalization
4. Gradient clipping

### Working Practically

"" While working with deep learning algorithms if you are facing a problem in which the loss from each or many number of epochs comes constant.

This is vanishing gradient problem as no an gradient is seen.

The above given methods can be used in such time to overcome the problem. ""

## 12. Hyperparameter Tuning

- 6 Hyperparameter tuning is a crucial step in optimizing deep learning models.
- 6 It involves finding the best set of hyperparameters that result in the most effective model.
- 6 It is used to increase accuracy.

## Key Hyperparameters

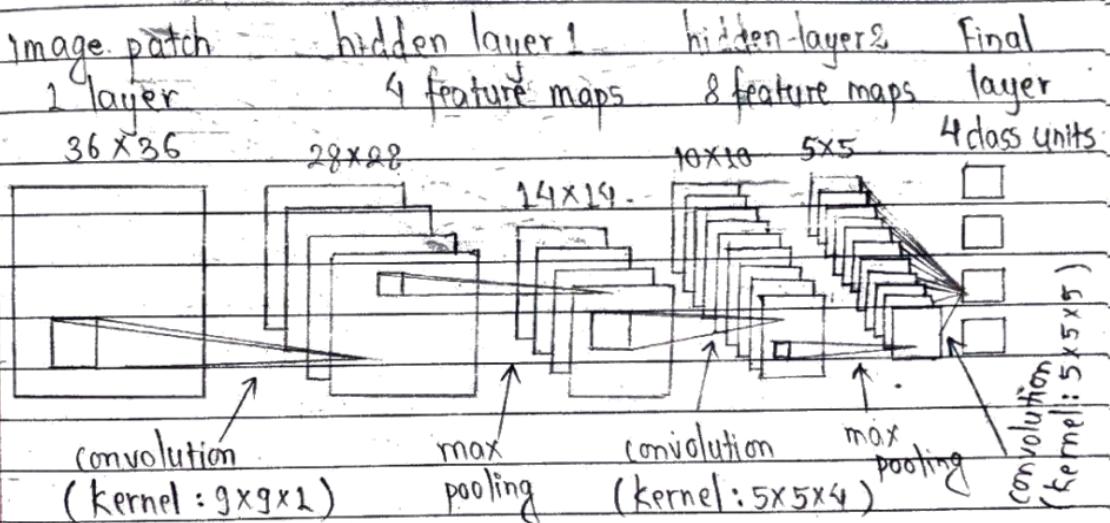
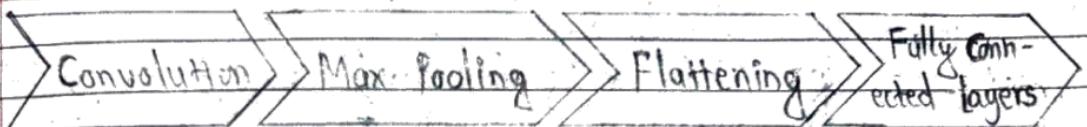
1. **Learning Rate:** Controls how much to change the model in response to the estimated error each time the model's weights are updated.
2. **Batch Size:** Number of training examples utilized in one iteration.
3. **Number of Epochs:** Number of complete passes through the training dataset.
4. **Number of Layers and Neurons:** Structure of the neural network, including the number of layers and the number of neurons in each layer.
5. **Activation Function:** Function that introduce non-linearity to the model.
6. **Optimizer:** Algorithm used to update the weights of network (e.g. Adam, SGD)
7. **Dropout Rate:** Fraction of input units to drop during training to prevent overfitting.
8. **Regularization Parameters:** Parameters like L2, L1 regularization to prevent overfitting.

## Methods for Hyperparameter Tuning

1. **Grid Search :** Exhaustively searches through a manually specified subset of the hyperparameter space.
2. **Random Search :** Randomly samples the hyperparameter space and performs better than grid search in practice.
3. **Bayesian Optimization :** Uses a probabilistic model to find the best hyperparameters.
4. **Hyperband :** Combines random search with early stopping to efficiently find the best hyperparameters.
5. **Genetic Algorithm :** Uses principles of natural selection to search for optimal hyperparameters.

### 13. Convolutional Neural Network: (CNN)

6. CNN is a type of artificial neural network, which is widely used for image/object recognition and classification.
7. Deep learning thus recognizes objects in an image by using a CNN.
8. This works like our eyes for capturing data and brain for predictions.



## 1 Convolution

Now in mathematics convolution is a mathematical operation on two functions that produces a third function that expresses how the shape of one is modified by the other.

Mathematically,

$$c = (f * g)(t) = \int_0^t f(\tau) g(t - \tau) d\tau$$

Where,

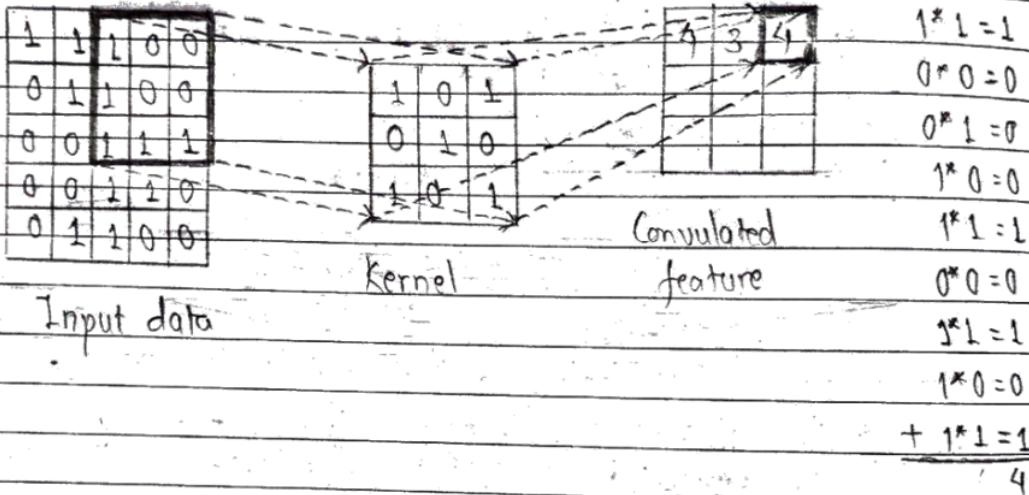
$c$  = convolution output function

$f$  = original output function

$g$  = function that is shifted over the input

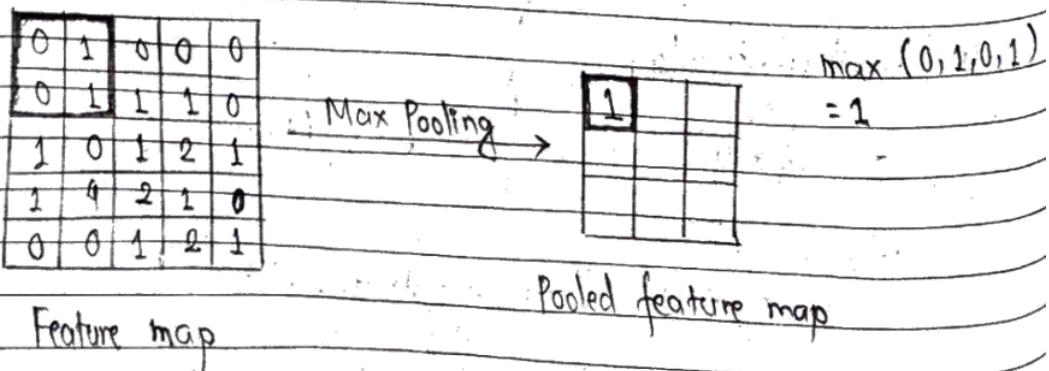
function.

$t$  = variable representing range of shifting  
 $\gamma$  = shifting against  $t$



## 2 Pooling

- 6 A pooling layer is another building block of a CNN.
- 6 Its function is to progressively reduce the spatial size of the representation to reduce the amount of parameters and computation in the network.



### 3. Flattening

- Flattening is converting the data into a 1-dimensional array for inputting it to next layer.
- Flatten output of the convolutional layers to create a single long feature vector.
- It is connected to the final classification model, which is called a fully connected layer.

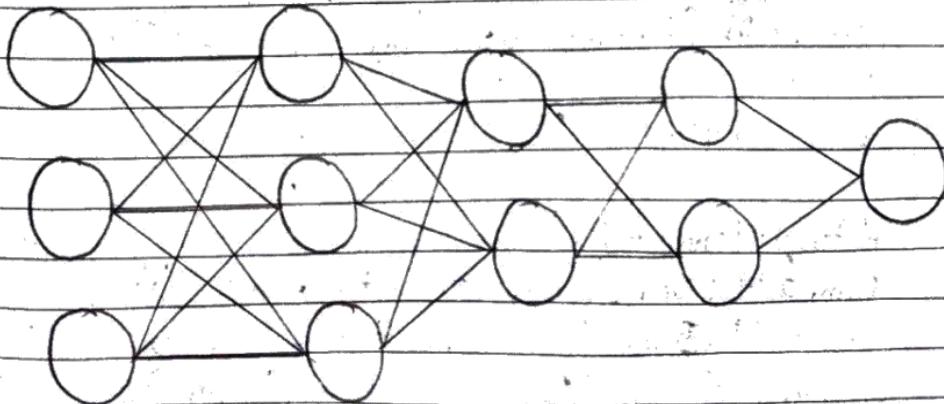
	$x_1$	$x_2$	$x_3$	$x_4$	$x_5$	$x_6$	$x_7$	$x_8$	$x_9$	$y$
1 1 0	1	1	0	4	2	1	0	2	1	-
4 2 1										
0 2 1										

flattening

pooled feature

### 4. Fully Connected Layers

- It is an Artificial Neural Network (ANN)



## Working Practically

```
import tensorflow  
from keras.layers import Dense, Conv2D, MaxPooling2D,  
                        flatten
```

```
from keras.models import Sequential
```

```
from tensorflow.keras.preprocessing.image import  
    ImageDataGenerator
```

```
# Create a CNN network.
```

```
cnn = Sequential()
```

```
# Convolution
```

```
cnn.add(Conv2D(filters=32, kernel_size=(3,3),  
               input_shape=(64, 64, 3),  
               activation="relu"))
```

```
# Pooling
```

```
cnn.add(MaxPooling2D(pool_size=(2,2)))
```

```
# Convolution
```

```
cnn.add(Conv2D(filters=32, kernel_size=(3,3),  
               activation="relu"))
```

```
# Pooling
```

```
cnn.add(MaxPooling2D(pool_size=(2,2)))
```

```
# flatten
```

```
cnn.add(Flatten())
```

```
# Fully Connected Layers
```

```
cnn.add(Dense(64, activation="relu"))
```

```
cnn.add(Dense(32, activation="relu"))
```

```
cnn.add(Dense(16, activation="relu"))
```

```
cnn.add(Dense(8, activation="relu"))
```

```
cnn.add(Dense(4, activation = "relu"))
cnn.add(Dense(2, activation = "relu"))
cnn.add(Dense(1, activation = "sigmoid"))
```

# Compile model

```
cnn.compile(optimizer = "adam", loss = "binary_crossentropy")
```

# Create an instance of ImageDataGenerator

```
train_datagen = ImageDataGenerator(
    rescale = 1.0 / 255,
    shear_range = 0.2,
    zoom_range = 0.2,
    horizontal_flip = True
)
```

# Same for test\_datagen

```
test_datagen = ImageDataGenerator(rescale = 1.0 / 255)
```

# Generate batches of augmented image data from  
# the specified directory

```
train_generator = train_datagen.flow_from_directory(
    ".\Pictures\Train",
    target_size = (64, 64),
    batch_size = 32,
    class_mode = "binary"
)
```

# Same for test\_generator

```
test_generator = test_datagen.flow_from_directory(
    ".\Pictures\Test",
    target_size = (64, 64),
    batch_size = 32,
    class_mode = "binary"
)
```

# Train the model using fit method

```
cnn.fit(
```

train\_generator,

steps\_per\_epoch=20,

epochs=5,

validation\_data=test\_generator)

```
)
```

# Evaluate the model on the test data

```
results=cnn.evaluate(test_generator)
```

results

# Making predictions (new)

```
from keras.preprocessing import Image
```

# Load image

```
img=Image.load_img(r".\Pictures\Predict\ Dogs\\PredictDog1.jpg",target_size=(64,64))
```

img

# Convert the image to array

```
img=Image.img_to_array(img)
```

img

# Change the dimensions of image

```
import numpy as np
```

```
img=np.expand_dims(img, axis=0)
```

cnn.predict(img) # New prediction

# To get the output in desired form

p = cnn.predict(img)

if p[0][0] < 0.5,

print("Dog")

else:

print("Cat")

""" Predict the multiple values from the folder """

# Create an Image Generator for prediction

predict\_datagen = ImageGenerator(rescale=1.0/225)

# Generate batches of images from the specified

# directory for prediction

predict\_generator = predict\_datagen.flow\_from\_directory(

r"\Pictures\Predict",

target\_size=(64, 64), batch\_size=32,

class\_mode=None,

shuffle=False

)

# Make predictions

predictions = cnn.predict(predict\_generator)

# Convert predictions to class labels

threshold = 0.5

predicted\_classes = np.where(predictions > threshold, "Dog", "Cat")

```
# Get the filenames
```

```
filenames = predict_generator.filenames
```

```
# Combine files and predictions
```

```
results = list(zip(filenames, predicted_classes))
```

```
# Print the results
```

```
for filename, prediction in results:
```

```
    print(f"file : {filename}, Predicted Class :  
          {prediction}")
```